Alex Rivera
Senior Machine Learning Engineer
San Francisco, CA | alex.rivera@example.com | github.com/arivera-ml

PROFESSIONAL SUMMARY
Results-driven Machine Learning Engineer with 6+ years of experience designing and deploying scalable deep learning syste

TECHNICAL SKILLS
- Languages: Python, C++, SQL, Bash
- Frameworks: PyTorch, Hugging Face Transformers, FastAPI, Ray, TensorRT
- Cloud & MLOps: AWS (SageMaker, EKS), Docker, Kubernetes, MLflow, Weights & Biases
- Data: PostgreSQL, Redis, FAISS, Pinecone, Apache Spark

WORK EXPERIENCE
Staff ML Engineer | DeepTech AI (2022 - Present)
- Architected semantic search engine serving 20M+ queries daily using fine-tuned BERT models.
- Reduced inference latency by 42% through INT8 quantization and ONNX runtime optimizations.
- Built automated evaluation pipelines tracking precision@k and NDCG across production deployments.

Senior AI Developer | Quantum Analytics (2019 - 2022)
- Developed bi-encoder and cross-encoder architectures for candidate-to-job recommendation.
- Led migration of batch training jobs to distributed PyTorch on multi-GPU clusters.

EDUCATION
M.S. in Computer Science (Machine Learning) | Stanford University (2019)
B.S. in Computer Engineering | UC Berkeley (2017)